

```

import cv2
import os
import random
import numpy as np
from matplotlib import pyplot as plt

from tensorflow.keras.models import Model
from tensorflow.keras.layers import Layer, Conv2D, Dense, MaxPooling2D, Input, Flatten

POS_PATH = "/content/sample_data/positive"
NEG_PATH = "/content/sample_data/negative"
ANC_PATH = "/content/sample_data/anchor"

# os.getcwd()
# /content/sample_data/anchor

import tensorflow as tf

anchor = tf.data.Dataset.list_files(ANC_PATH+'/*.jpg').take(300)
positive = tf.data.Dataset.list_files(POS_PATH+'/*.jpg').take(300)
negative = tf.data.Dataset.list_files(NEG_PATH+'/*.jpg').take(300)

def preprocess(file_path):

    byte_img = tf.io.read_file(file_path)
    img = tf.io.decode_jpeg(byte_img)
    img = tf.image.resize(img, (100,100))
    img = img / 255.0

    return img

positives = tf.data.Dataset.zip((anchor, positive, tf.data.Dataset.from_tensor_slices(tf.ones(len(anchor)))))
negatives = tf.data.Dataset.zip((anchor, negative, tf.data.Dataset.from_tensor_slices(tf.zeros(len(anchor)))))
data = positives.concatenate(negatives)

sample = data.as_numpy_iterator()

eg = sample.next()
eg

(b'/content/sample_data/anchor/461bcb07-2bb9-11ee-ac0d-204ef6489e9c.jpg',
 b'/content/sample_data/positive/810acd42-2bba-11ee-b148-204ef6489e9c.jpg',
 1.0)

def preprocess_twin(input_img, validation_img, label):
    return (preprocess(input_img) , preprocess(validation_img) , label)

# preprocess_twin(*eg)
res = preprocess_twin(*eg)
plt.imshow(res[0])

```

```
<matplotlib.image.AxesImage at 0x7a4724hd7190>
data = data.map(preprocess_twin)
data = data.cache()
data = data.shuffle(buffer_size = 1024)
```




```
# training partition
train_data = data.take(round(len(data)*0.7))
train_data = train_data.batch(16)
train_data = train_data.prefetch(8)
```



```
train_samples = train_data.as_numpy_iterator()
train_sample = train_samples.next()
plt.imshow(train_sample[0][6])
```

```
<matplotlib.image.AxesImage at 0x7a47200ee710>
```



```
# we have batch of 16
len(train_sample[0])
```

```
16
```

```
# testing partition
test_data = data.skip(round(len(data)*0.7))
test_data = test_data.batch(16)
test_data = test_data.prefetch(8)
# test_data = data.take(round(len(data)*0.3)) # no sense writing it
```

```
def make_embedding():
```

```
# first block
inp = Input(shape = (100, 100, 3), name = 'input_image')
c1 = Conv2D(64, (10, 10), activation = 'relu')(inp)
m1 = MaxPooling2D(64, (2, 2), padding = 'same')(c1)

# second block
c2 = Conv2D(128, (7,7), activation = 'relu')(m1)
m2 = MaxPooling2D(64, (2, 2), padding = 'same')(c2)

# third block
c3 = Conv2D(128, (4,4), activation = 'relu')(m2)
m3 = MaxPooling2D(64, (2, 2), padding = 'same')(c3)

# fourth layer
c4 = Conv2D(256, (4,4), activation = 'relu')(m3)
f1 = Flatten()(c4)
d1 = Dense(4096, activation = 'sigmoid')(f1)

return Model(inputs = [inp], outputs = [d1], name = 'embedding')
```

```
embedding = make_embedding()
embedding.summary()
```

```
Model: "embedding"
```

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```
=====
input_image (InputLayer)      [(None, 100, 100, 3)]      0

conv2d (Conv2D)                (None, 91, 91, 64)      19264

max_pooling2d (MaxPooling2D   (None, 46, 46, 64)      0
)

conv2d_1 (Conv2D)              (None, 40, 40, 128)     401536

max_pooling2d_1 (MaxPooling   (None, 20, 20, 128)     0
2D)

conv2d_2 (Conv2D)              (None, 17, 17, 128)     262272

max_pooling2d_2 (MaxPooling   (None, 9, 9, 128)       0
2D)

conv2d_3 (Conv2D)              (None, 6, 6, 256)       524544

flatten (Flatten)              (None, 9216)            0

dense (Dense)                  (None, 4096)            37752832

=====
Total params: 38,960,448
Trainable params: 38,960,448
Non-trainable params: 0
=====

# will create a custom nn layer
class L1Dist(Layer):

    def __init__(self, **kwargs):
        super().__init__()

    # the crux of the paper lies here(charaterstics of the siamese nn)!
    def call(self, input_embedding, validation_embedding):
        return tf.math.abs(input_embedding - validation_embedding)

l1 = L1Dist()
l1

<__main__.L1Dist at 0x7a46eb9a0e20>

def make_siamese_model():

    input_image = Input(name = 'input_img', shape = (100, 100, 3))
    validation_image = Input(name = 'validation_img', shape = (100, 100, 3))

    siamese_layer = L1Dist(name = 'distance')
    # siamese_layer.name = 'distance'
    distances = siamese_layer(embedding(input_image), embedding(validation_image))

    # final classification layer
    classifier = Dense(1, activation = 'sigmoid')(distances)

    return Model(inputs = [input_image, validation_image], outputs = classifier, name = 'SiameseNetwork')

siamese_model = make_siamese_model()
siamese_model.summary()

Model: "SiameseNetwork"
-----
Layer (type)                Output Shape                Param #   Connected to
-----
input_img (InputLayer)      [(None, 100, 100, 3  0    []
)]

validation_img (InputLayer) [(None, 100, 100, 3  0    []
)]

embedding (Functional)      (None, 4096)               38960448  ['input_img[0][0]',
                                'validation_img[0][0]']

l1_dist_1 (L1Dist)          (None, 4096)               0         ['embedding[0][0]',
                                'embedding[1][0]']

dense_1 (Dense)             (None, 1)                  4097      ['l1_dist_1[0][0]']

=====
Total params: 38,964,545
```

Trainable params: 38,964,545
 Non-trainable params: 0

```
binary_cross_loss = tf.losses.BinaryCrossentropy() # recommended set logits = true when the data is not standerizer
opt = tf.keras.optimizers.Adam(1e-4)
```

```
checkpoint_dir = "./training_checkpoints"
checkpoint_prefix = os.path.join(checkpoint_dir, 'ckpt')
checkpoints = tf.train.Checkpoint(opt = opt, siamese_model = siamese_model)
```

```
test_batch = test_data.as_numpy_iterator().next()
len(test_batch)
```

```
# understanding the data
```

```
# np.array(test_batch).shape
# x = test_batch[:2]
# np.array(x).shape
# len(test_batch[1])
```

```
# np.array(sample[0]).shape
# len(sample[0][0][0][0])
```

```
3
```

```
@tf.function # this line will make the follwing function will integrate the following function to be a computational grpah!
def train_step(batch):
```

```
    # recording all of our operations
    with tf.GradientTape() as tape:
```

```
        # getting anchor and positive images
        X = batch[:2]
        y = batch[2]
```

```
        # forward pass
        yhat = siamese_model(X, training = True) # setting training = true madetory for training
        # calculate loss
        loss = binary_cross_loss(y, yhat)
```

```
    print(loss)
```

```
    # calculating gradients
    grad = tape.gradient(loss, siamese_model.trainable_variables)
```

```
    # calcuates updated weights and apply to the siamese network
    opt.apply_gradients(zip(grad, siamese_model.trainable_variables))
```

```
    return loss
```

```
def train(data, EPOCHS):
```

```
    for epoch in range(1, EPOCHS + 1):
```

```
        print('\n Epoch {}/{}'.format(epoch, EPOCHS))
        progbar = tf.keras.utils.Progbar(len(data)) # to inculcate the progression!
```

```
        for idx, batch in enumerate(data):
            train_step(batch)
            progbar.update(idx + 1)
```

```
        if epoch % 10 == 0:
            checkpoints.save(file_prefix = checkpoint_prefix)
```

```
EPOCHS = 50
train(train_data, EPOCHS)
```

```

Epoch 35/50
27/27 [=====] - 5s 190ms/step

Epoch 36/50
27/27 [=====] - 5s 190ms/step

Epoch 37/50
27/27 [=====] - 5s 190ms/step

Epoch 38/50
27/27 [=====] - 5s 190ms/step

Epoch 39/50
27/27 [=====] - 5s 191ms/step

Epoch 40/50
27/27 [=====] - 5s 191ms/step

Epoch 41/50
27/27 [=====] - 5s 190ms/step

Epoch 42/50
27/27 [=====] - 5s 190ms/step

Epoch 43/50
27/27 [=====] - 5s 190ms/step

Epoch 44/50
27/27 [=====] - 5s 190ms/step

Epoch 45/50
27/27 [=====] - 5s 191ms/step

Epoch 46/50
27/27 [=====] - 5s 190ms/step

Epoch 47/50
27/27 [=====] - 5s 190ms/step

Epoch 48/50
27/27 [=====] - 5s 190ms/step

Epoch 49/50
27/27 [=====] - 5s 190ms/step

Epoch 50/50
27/27 [=====] - 5s 190ms/step

from tensorflow.keras.metrics import Precision, Recall

test_input, test_val, y_true = test_data.as_numpy_iterator().next()

test_var = test_data.as_numpy_iterator().next()
test_var

len(test_input[0])

100

siamese_model.save('siamese_model.h5')

the loaded model, but the compiled metrics have yet to be built. `model.compiled_metrics` will be empty until you train or evaluate

# siamese_model.predict(test_input, test_val)

siamese_model.predict([test_input, test_val])

y_hat = siamese_model.predict([test_input, test_val])
y_hat

# post processing the results

[1 if prediction > 0.5 else 0 for prediction in y_hat]

[0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0]

y_true

```

```
array([0., 0., 0., 1., 0., 1., 0., 0., 1., 0., 1., 1., 0., 0., 0., 0.],
      dtype=float32)
```

```
m = Recall()
m.update_state(y_true, y_hat)
m.result().numpy()
```

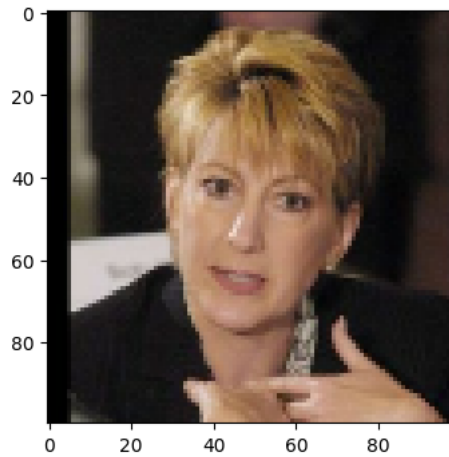

```
1.0
```

```
m = Precision()
m.update_state(y_true, y_hat)
m.result().numpy()
```



```
1.0
```

▼ visulatizing the results



```
plt.figure(figsize=(10,4))
plt.subplot(1, 2, 1)
plt.imshow(test_input[0])
plt.subplot(1, 2, 2)
plt.imshow(test_val[0])
plt.show()
```



```
plt.figure(figsize=(10,4))
plt.subplot(1, 2, 1)
plt.imshow(test_input[3])
plt.subplot(1, 2, 2)
plt.imshow(test_val[3])
plt.show()
```



```
plt.figure(figsize=(10,4))
plt.subplot(1, 2, 1)
plt.imshow(test_input[5])
plt.subplot(1, 2, 2)
plt.imshow(test_val[5])
plt.show()
```



Double-click (or enter) to edit

●